

Fitting spatial and spatiotemporal models - continued

FISH 507 – Applied Time Series Analysis

Eric Ward

4 Mar 2021

Where we left off

- ▶ Many approaches we've talked about were forms of spatial regression

$$Y_i = \mathbf{b}X_i + \phi_i + \varepsilon_i$$

- ▶ ϕ is a vector of random effects describing a spatial field / process
- ▶ Is this a state space model?

Gaussian predictive process model

- ▶ Instead of estimating giant covariance matrix describing ϕ ,
- ▶ Estimate smaller random effects at a subset of points ϕ^*
- ▶ Use some covariance kernel, and distances between points describing ϕ and ϕ^* to project estimates
- ▶ Can be done in ML or Bayesian framework, Latimer et al. 2009

Method 3: glmmfields

- ▶ 'glmmfields' implements these models using Stan backend
- ▶ Easy to use formula syntax like `gam()` or `spBayes()`
- ▶ Allows spatial field to be modeled with Multivariate-T field, as alternative to Multivariate-Normal to better capture extremes
- ▶ Also includes lots of non-normal families for observation model

glmmfields

Let's start with a simple model. This includes only 6 knots (probably way too few) and 100 iterations (too few!) for run time.

- ▶ Time set to NULL here to fit a model with constant spatial field

```
m <- glmmfields(Feb ~ 0, time = NULL,  
  lat = "Latitude", lon = "Longitude", data = d,  
  nknots = 6, iter = 100, chains = 2, seed = 1)
```

glmmfields

Ok, now let's change the covariance function to matern, and estimate separate spatial fields by year.

- ▶ The 'estimate_ar' argument is important for determining whether the field is an AR(1) process or independent by year

```
m <- glmmfields(Feb ~ 0, time = "Water.Year",  
  lat = "Latitude", lon = "Longitude", covariance="matern",  
  data = d, nknots = 6, iter = 100, chains = 2, seed = 1,  
  estimate_ar=FALSE)
```

glmmfields

Finally let's include an example with modeling the spatial field as a Multivariate-t distribution

- We do this with the 'estimate_df' argument

```
m <- glmmfields(Feb ~ 0, time = "Water.Year",  
  lat = "Latitude", lon = "Longitude", covariance="matern",  
  data = d, nknots = 6, iter = 100, chains = 2, seed = 1,  
  estimate_AR=FALSE, estimate_df= TRUE)
```

Method 4: INLA

GP spatial models like glmmfields are extremely powerful

- ▶ May get overwhelmed by number of points
- ▶ Approaches to incorporate knots using sparse covariance matrices
- ▶ Integrated Nested Laplace Approximation - not on CRAN

Motivation of INLA

- ▶ Some problems contain simple spatial structure
 - ▶ e.g. the harbor seal data in MARSS() with only 5 – 7 time series
- ▶ Others are much more complex
 - ▶ WA SWE data
 - ▶ Fisheries survey data (1000s of points)
- ▶ Including time-varying spatial fields becomes very computationally difficult
- ▶ Doing all of the above in a Bayesian setting can be prohibitive, but we can use Laplace approximation

INLA's approximation: SNOTEL data

- ▶ INLA meshes describe resolution of surface to estimate spatial process

```
INLA::meshbuilder
```

INLA's approximation: SNOTEL data

- ▶ How many points fall on vertices?
- ▶ Is the boundary area large enough?
- ▶ Choosing this must be done very carefully!

Estimation done via maximum likelihood

- ▶ Estimates seems similar to those from `gls()` and `spBayes()`
- ▶ Year included as numeric here (not significant)
- ▶ Alternatively, we can include year in spatial field
- ▶ Year can also be included as factor
- ▶ splines, AR coefficients also possible for time-varying effects


Projecting INLA estimates to surface

Spatial SWE fields by year

Implementation in sdmTMB

- Start with the built in `make_mesh` function

```
snow_spde <- make_mesh(d, c("Longitude", "Latitude"), n_knn = 5)  
plot(snow_spde)
```



Estimation in INLA

- ▶ Can be slow for very large models
 - ▶ but can also generate approximate Bayesian estimates
- ▶ Alternatively using automatic differentiation / TMB to maximize likelihood
 - ▶ VAST
 - ▶ sdmTMB

Implementation in sdmTMB

- ▶ This would fit a model with static spatial field

```
fit <- sdmTMB(Feb ~ 1, spde = snow_spde, data=d)
```

- ▶ Versus this one, where the field is independent by year


```
fit <- sdmTMB(Feb ~ 1, spde = snow_spde, time="Water.Year")
```

Diagnostics: residuals for a given time slice

```
fit <- sdmTMB(Feb ~ 1, spde = snow_spde, data=d)
qqnorm(residuals(fit))
```



Diagnostics: residuals over time: are they stationary?



INLA & TMB references:

Lindgren 2013

Lindgren & Rue 2015

INLA book

Kristensen et al. 2016